













ggplot2



First Things First: Course Evaluation



<https://uzk-evaluation.uni-koeln.de/evasys/online.php?p=W4L6C>

Recap: Data Visualization

- 1) Why use data visualization?
- 2) What makes a good data visualization?
- 3) Think about your audience – Who?
- 4) Tables, charts and colors
- 5) Common “zero-code” tools for data visualization




If you have not downloaded the files for today yet:

- taxation.csv
- sales.csv
- ggplot_intro.r

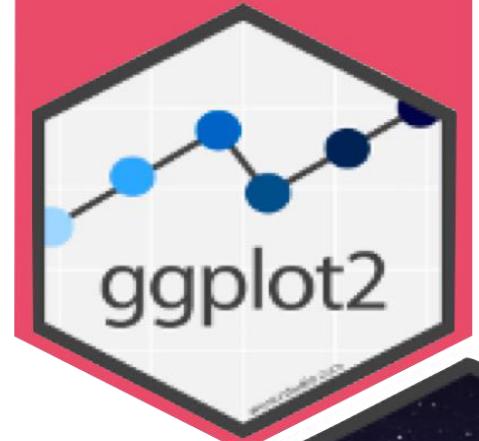
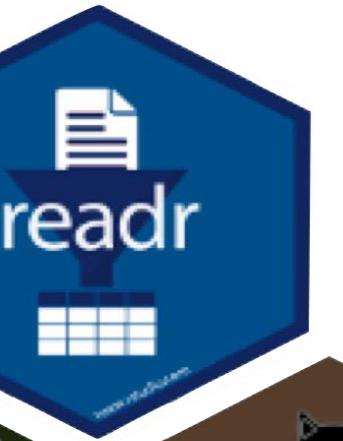
Please do so now.

Agenda

- ggplot
- ggplot practice




Intro to ggplot

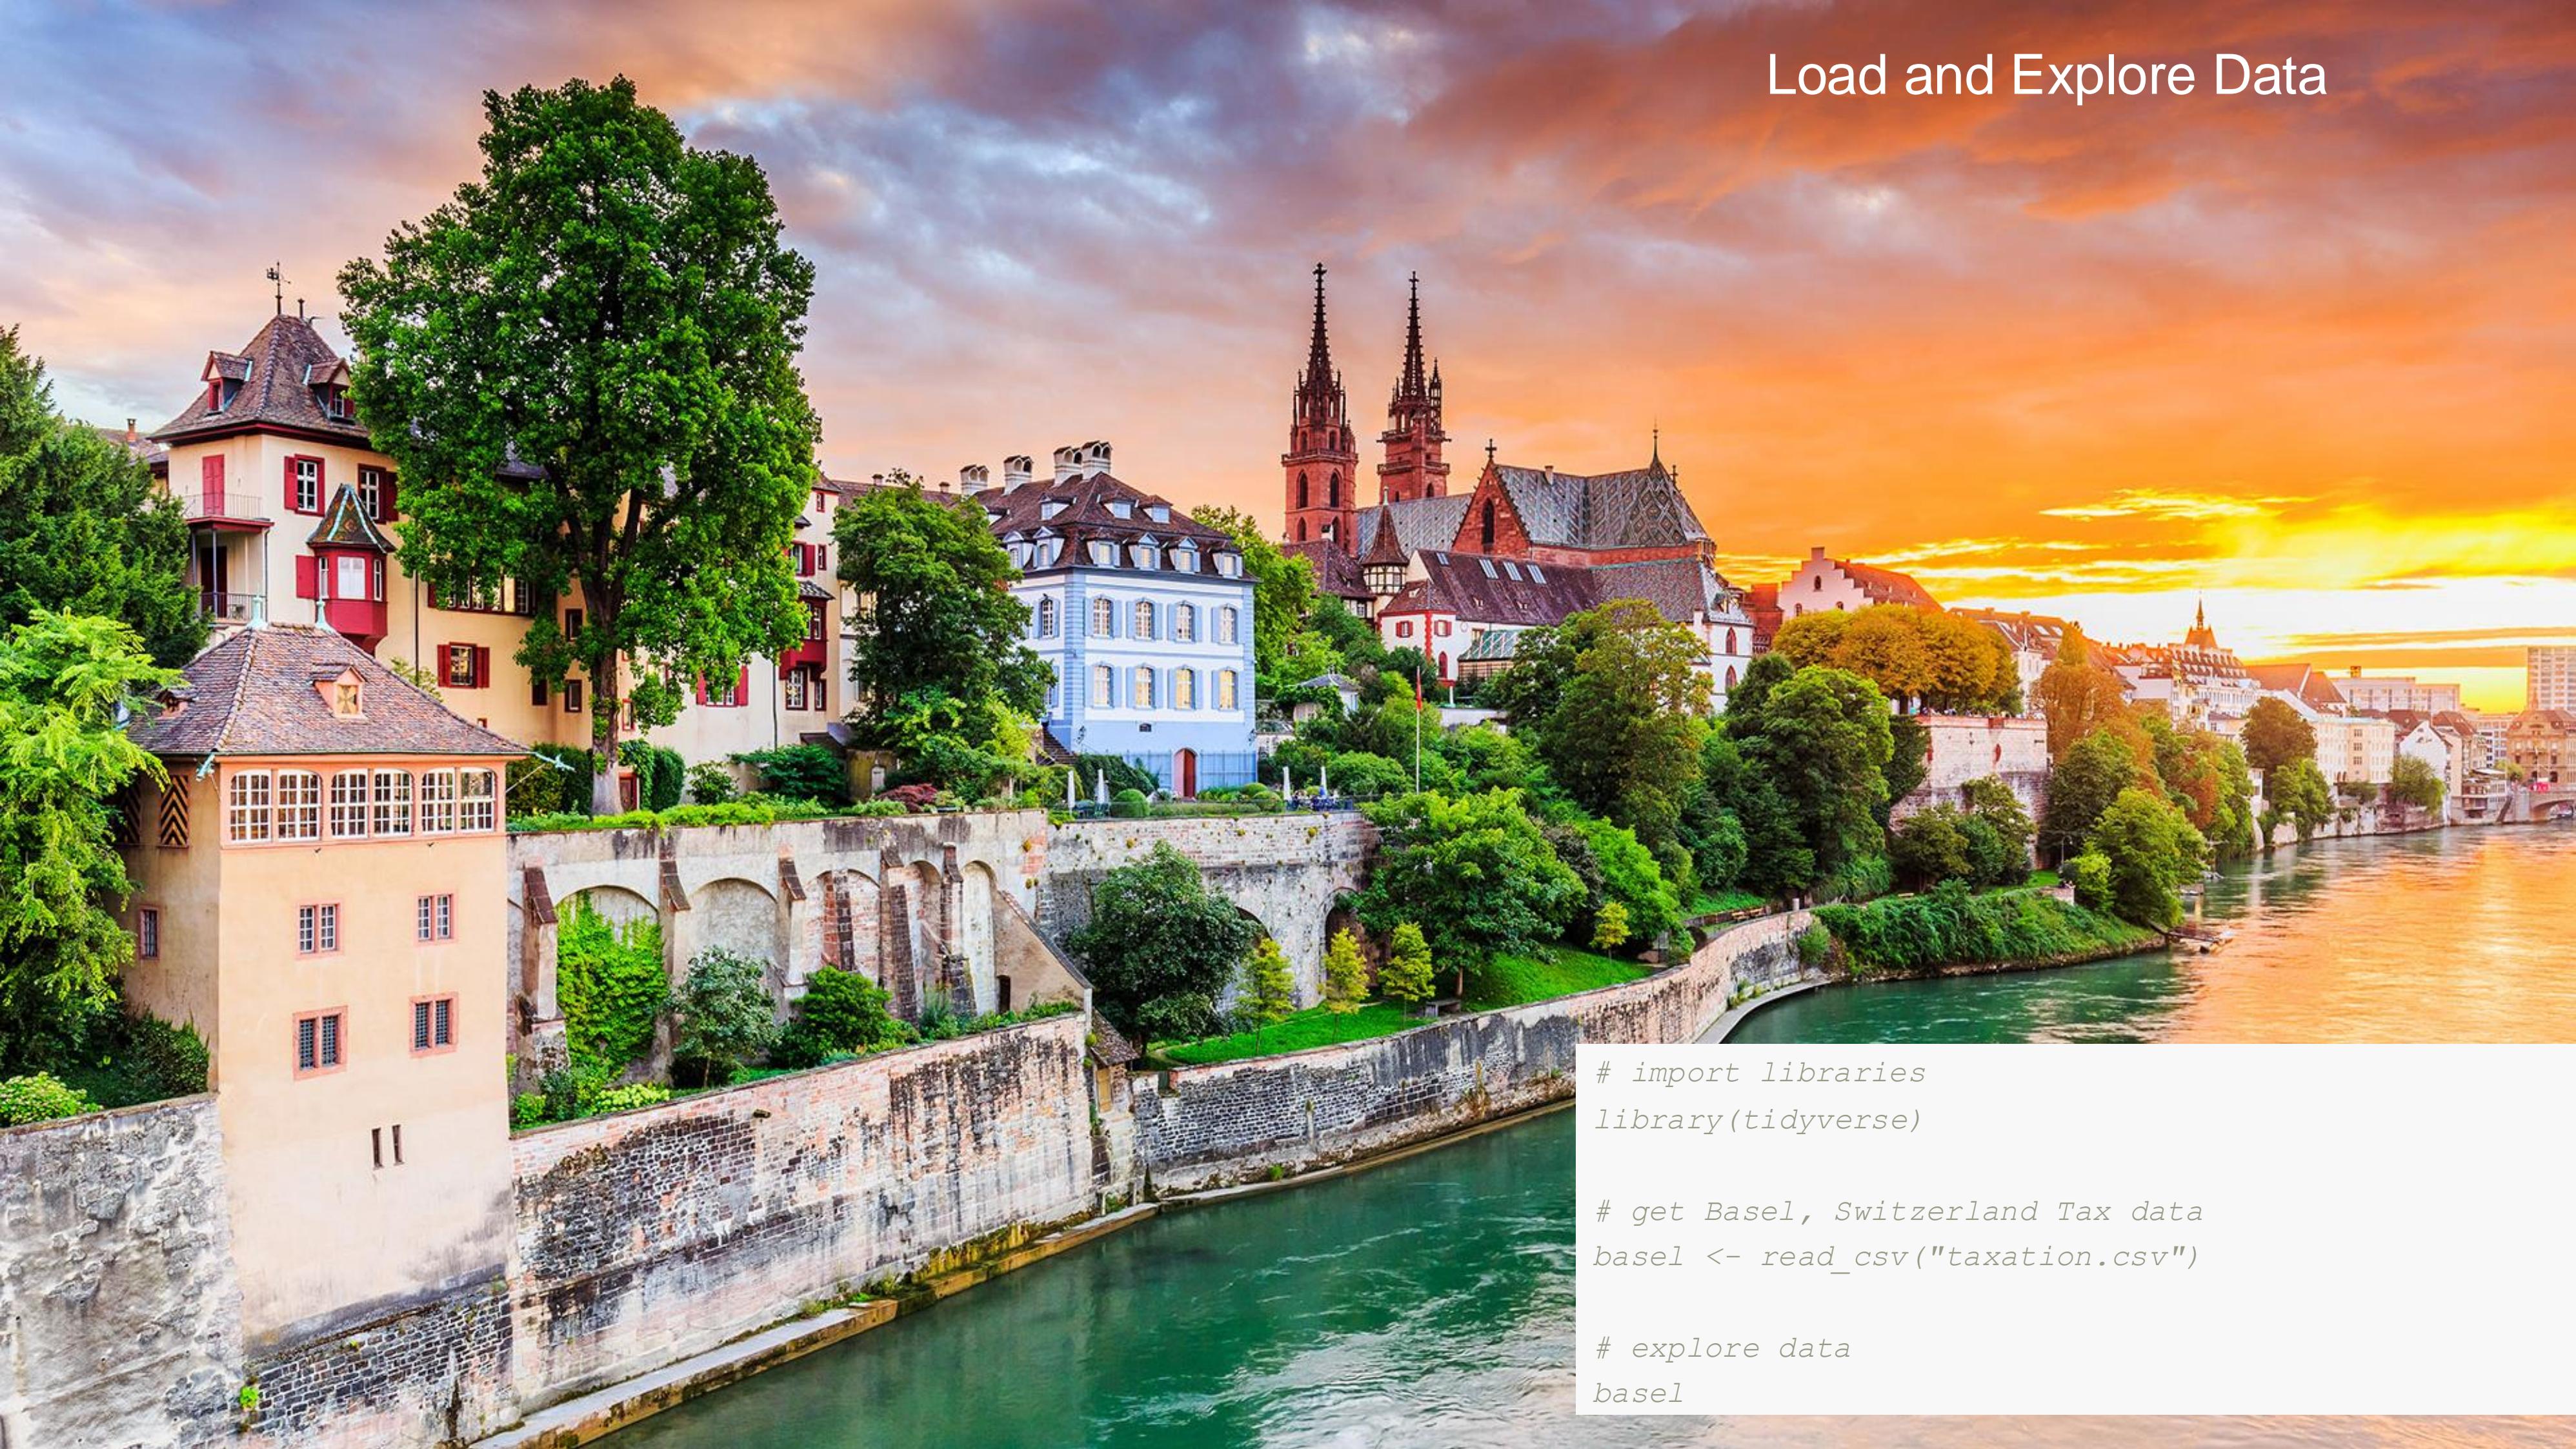
ggplot2 Visualization for exploration and reporting 	dplyr Data organization, transformation, and summaries 	stringr Text analysis using regular expressions 	readr Reading rectangular text files (.txt, .csv) 
tibble Lean version of the data frame 	tidyverse A collection of R packages for data science 	purrr Functional programming 	forcats Handling factors 

ggplot2 uses layers to create plots

- 1) **data**: the data set
- 2) **mapping**: the plot's structure
 - What do the axes represent?
 - Map variables to colors, shapes, sizes
- 3) **geoms**: geometric shapes illustrating data
- 4) **facets**: stratify the plot according to variable
- 5) **labs**: plot annotation
- 6) **themes**: aesthetic details
- 7) **scales**: scaling of dimensions



Load and Explore Data



```
# import libraries  
library(tidyverse)  
  
# get Basel, Switzerland Tax data  
basel <- read_csv("taxation.csv")  
  
# explore data  
basel
```

Data : taxation.csv

```
# A tibble: 357 × 10
  year quarter    quarter_no     N income_mean income_median income_gini wealth_mean wealth_median wealth_gini
  <dbl> <chr>        <dbl> <dbl>       <dbl>       <dbl>       <dbl>       <dbl>       <dbl>
1 2001 Altstadt Grossbasel      1 1673     87776      51819      0.593     1014720     20201      0.952
2 2001 Vorstädte             2 3204     84109      49914      0.577     1119418     19045      0.957
3 2001 Am Ring              3 6579     62582      49426      0.467     300878      16024      0.879
4 2001 Breite                4 5433     52039      47227      0.358     105198      10820      0.826
5 2001 St. Alban             5 6179     89956      58112      0.54      778475      40315      0.901
6 2001 Gundeldingen          6 11224     51229      46265      0.387     92099       3437      0.871
7 2001 Bruderholz            7 5090     96124      64512      0.52      982401      63530      0.9
8 2001 Bachletten             8 8157     70348      56258      0.444     346088      32129      0.853
9 2001 Gotthelf               9 4256     59049      47960      0.435     324687      16650      0.916
10 2001 Iselin                 10 9853     49631      45530      0.371     99290       9065      0.832
# i 347 more rows
# i Use `print(n = ...)` to see more rows
```

Prepare dataset

Calculate mean + median income for each year

```
basel_avg <- basel %>%
  group_by(year) %>%
  summarize(income_mean = mean(income_mean),
            income_median = mean(income_median))
```

basel_avg

A tibble: 17 × 3

year	income_mean	income_median
1	63027.	49516.
2	63555.	50066.
3	63083.	49717.
4	62298.	49467.
5	63133.	49192.
6	64148.	49102.
7	66594.	50164.
8	66463.	48068.
9	66614.	48818.
10	67185.	49028.
11	66050.	49213.
12	66987.	49433.
13	68748.	49878.
14	70499.	50440.
15	71115.	50426.
16	73272.	50653.
17	72388.	50840.

Data

All plots start with `ggplot()`

```
ggplot(data = basel_avg)
```

Two arguments:

`data` | The data set (tibble)

`mapping` | The plot structure. Defined using `aes()`

aes()

aes() helps define the structure of the **mapping** argument.

Key arguments:

x, y | defines axes


color, fill | defines colors

alpha | defines opacity

size | defines sizes

shape | defines shapes (e.g., circles or squares)


```
ggplot(data = basel_avg,  
       mapping = aes(x = year,  
                      y = income_mean))
```



geom_*()

The + operator "adds" **additional elements** to the plot.

```
ggplot(data = basel_avg,  
       mapping = aes(x = year,  
                      y = income_mean)) +  
  
  # Show as points  
  geom_point()
```



geom_*()

geom_*() functions define which geometric objects are used to illustrate the data.

A few examples geoms:

geom_point() | for points


geom_line() | for lines

geom_smooth() | for smooth curves

geom_bar() | for bars

geom_boxplot() | for box-plots


geom_violin() | for violin-plots



geom_*()

geom_*() functions define which geometric objects are used to illustrate the data.


```
ggplot(data = basel_avg,  
       mapping = aes(x = year,  
                      y = income_mean)) +  
  
  # Show as lines  
  geom_line()
```



geom_*()

geom_*() functions define which geometric objects are used to illustrate the data.


```
ggplot(data = basel_avg,  
       mapping = aes(x = year,  
                      y = income_mean)) +  
  
  # Show as smoothed curve  
  geom_smooth()
```



geom_*()

geom_*() functions define which geometric objects are used to illustrate the data.


```
ggplot(data = basel_avg,  
       mapping = aes(x = year,  
                      y = income_mean)) +  
  
  # Show as points and lines  
  geom_point() +  
  geom_line()
```



geom_*()

geom_*() functions define which geometric objects are used to illustrate the data.

```
ggplot(data = basel_avg,  
       mapping = aes(x = year,  
                      y = income_mean)) +  
  
  # Add bars (not necessarily recommended)  
  geom_bar(stat = "identity") +  
  
  # Show as points and lines  
  geom_point() +  
  geom_line()
```




Colors

R understands a large number of **color names** (see `colors()` for the whole set).

Additionally colors can be specified using **hex codes** or the `rgb()` function.

```
ggplot(data = basel_avg,  
       mapping = aes(x = year,  
                      y = income_mean)) +  
  
  # Add bars (not necessarily recommended)  
  geom_bar(stat = "identity",  
            col = "lightblue",  
            fill = "lightblue") +  
  
  # Show as points and lines  
  geom_point(col = "#4682B4") +  
  geom_line(col = "#4682B4")
```




geom_*()

Most geom_*() functions allow the independent specification of **data** and **mapping**.

Can be used to add geoms for other cases or variables in the data.

```
ggplot(data = basel_avg,  
       mapping = aes(x = year,  
                      y = income_mean)) +  
  geom_point() +  
  geom_line() +  
  
  # Add points and lines for median  
  geom_point(aes(y = income_median)) +  
  geom_line(aes(y = income_median))
```



Wrangling

Oftentimes, creating the desired plot requires appropriate data wrangling.

ggplot works best with **long data formats**.


```
# pivot to long format
basel_avg_long <- basel_avg %>%
  pivot_longer(-year,
    names_to = "statistic",
    values_to = "income")
```

```
basel_avg_long
# A tibble: 34 × 3
  year statistic income
  <dbl> <chr>     <dbl>
1 2001 income_mean 63027.
2 2001 income_median 49516.
3 2002 income_mean 63555.
4 2002 income_median 50066.
5 2003 income_mean 63083.
6 2003 income_median 49717.
7 2004 income_mean 62298.
8 2004 income_median 49467.
9 2005 income_mean 63133.
10 2005 income_median 49192.
# ... with
```

aes()

aes() helps define the structure of the **mapping** argument


```
# use basel_avg_long
ggplot(data = basel_avg_long,
       mapping = aes(
         x = year,
         y = income,
         # add color dimension
         col = statistic)) +
  geom_point() +
  geom_line()
```



aes()

aes() helps define the structure of the **mapping** argument

```
# use basel_avg_long
ggplot(data = basel_avg_long,
       mapping = aes(
         x = year,
         y = income,
         # add shape dimension
         shape = statistic)) +
  geom_point() +
  geom_line()
```




aes()

aes() helps define the structure of the **mapping** argument

```
# use basel_avg_long
ggplot(data = basel_avg_long,
       mapping = aes(
         x = year,
         y = income,


         # add many dimensions
         # (not recommended)
         col = statistic,
         shape = statistic,
         size = statistic,
         alpha = statistic)) +
geom_point() +
geom_line()
```



aes()

aes() helps define the structure of the **mapping** argument

```
# use basel_avg_long
ggplot(data = basel_avg_long,
       mapping = aes(
         x = year,
         y = income,
         # add many dimensions
         col = statistic)) +
  geom_point() +
  geom_line()
```



facet_*()


Facetting creates the **same plot for groups** defined by another variable.

Key functions:

facet_wrap()

facet_grid()


```
basel_long <- basel %>%
  pivot_longer(c(income_mean, income_median),
               names_to = 'statistic',
               values_to = 'income')
```



facet_*()

Facetting creates the **same plot for groups** defined by another variable.

```
# use basel_long
ggplot(data = basel_long,
       mapping = aes(
         x = year,
         y = income,
         col = statistic)) +
  geom_point() +
  geom_line() +
  # facet by quarter
  facet_wrap(~quarter)
```



patchwork


patchwork provides a simple syntax to combine plots.

patchwork syntax:

- + | combine horizontally
- / | combine vertically
- | | spacer
- () | grouper
- & | apply to all
- plot_layout | control layout

```
# use patchwork
graph1 <- basel %>%
  filter(quarter == "Breite") %>%
  ggplot(aes(x = year,
             y = income_mean
            )) +
  geom_smooth()

graph2 <- basel %>%
  filter(quarter == "Bruderholz") %>%
  ggplot(aes(x = year,
             y = income_mean
            )) +
  geom_smooth()
```



patchwork

patchwork provides a simple syntax to combine plots.


patchwork syntax:

- + | combine horizontally
- / | combine vertically
- | | spacer
- () | grouper
- & | apply to all
- plot_layout | control layout

```
# use patchwork
graph1 <- basel %>%
  filter(quarter == "Breite") %>%
  ggplot(aes(x = year,
             y = income_mean
            )) +
  geom_smooth()

graph2 <- basel %>%
  filter(quarter == "Bruderholz") %>%
  ggplot(aes(x = year,
             y = income_mean
            )) +
  geom_smooth()
```

graph1+graph2



income_mean

120000
100000
80000

2005 2010 2015

year

year

income_mean

120000
100000
80000

2005 2010 2015

year

patchwork


patchwork provides a simple syntax to combine plots.

patchwork syntax:

- + | combine horizontally
- / | combine vertically
- | | spacer
- () | grouper
- & | apply to all
- plot_layout | control layout

```
# use patchwork
graph1 <- basel %>%
  filter(quarter == "Breite") %>%
  ggplot(aes(x = year,
             y = income_mean
            )) +
  geom_smooth()

graph2 <- basel %>%
  filter(quarter == "Bruderholz") %>%
  ggplot(aes(x = year,
             y = income_mean
            )) +
  geom_smooth()
```




patchwork

```
# use patchwork
graph1 <- basel %>%
  filter(quarter == "Breite") %>%
  ggplot(aes(x = year,
             y = income_mean
           )) +
  geom_smooth()

graph2 <- basel %>%
  filter(quarter == "Bruderholz") %>%
  ggplot(aes(x = year,
             y = income_mean
           )) +
  geom_smooth()

all <- basel %>%
  ggplot(aes(x = year,
             y = income_mean
           )) +
  geom_smooth()
```


```
all / (graph1+graph2) +
  plot_layout(guides="collect")
```




annotation

```
# Annotation example

basel %>%
  filter(quarter == "Bruderholz") %>%
  ggplot(aes(x = year,
             y = income_mean
  )) +
  annotate(geom="text", x=2010, y=140000,
           label="This is a big spike!") +
  annotate(geom="point",
           x=2010, y=137000, size=10,
           shape=21, fill="transparent") +
  geom_line()
```



Practice



```
# clean it up a bit  
# change format of 'Order Date' to "%m/%d/%y"  
# Use "%Y" to create a new variable "Year"  
# Define "Year" as numeric
```

```
# Create a graph that shows a point for the total profit for each year  
# Connect the dots using a line.
```

Add the total sales as well to the graph

```
# make the profit darkgreen
```

```
# change the shape of the profit to 0 (square)
```

```
# create a bar chart summarizing total sales by segment
```

```
# assign the color of the "fill" by the total sales
```

```
# create a bar chart that shows the sales sum by state instead  
# of segment. Color the fill again by total sales. Flip the| coordinates  
# using coord_flip()
```

```
# let's reorders state by Sales
```

```
# let's only take the top 10
```

```
# create a plot of the total profit for each region.  
Use facet_wrap()
```

```
# create a line chart of total profit of Furniture, Office Supplies,  
# and Everything (no filter) by year. Show the overall summary up top and the  
# furniture and office supplies side-by-side underneath
```

```
# so a line graph of mean sales across years. On the value for 2013  
# provide an annotation that we had a "Successful Social Media Campaign"  
# that year
```

Make a graph using something other than lines and dots (see slides)

`geom_abline()` `geom_hline()` `geom_vline()`
Reference lines: horizontal, vertical, and diagonal



`geom_bar()` `geom_col()` `stat_count()`
Bar charts



`geom_bin_2d()` `stat_bin_2d()`
Heatmap of 2d bin counts



`geom_blank()`
Draw nothing



`geom_boxplot()` `stat_boxplot()`
A box and whiskers plot (in the style of Tukey)



`geom_contour()` `geom_contour_filled()` `stat_contour()`
`stat_contour_filled()`
2D contours of a 3D surface



`geom_count()` `stat_sum()`
Count overlapping points



`geom_density()` `stat_density()`
Smoothed density estimates



`geom_density_2d()` `geom_density_2d_filled()` `stat_density_2d()`
`stat_density_2d_filled()`
Contours of a 2D density estimate



`geom_dotplot()`
Dot plot



`geom_errorbarh()`
Horizontal error bars



`geom_function()` `stat_function()`
Draw a function as a continuous curve



`geom_hex()` `stat_bin_hex()`
Hexagonal heatmap of 2d bin counts



`geom_freqpoly()` `geom_histogram()` `stat_bin()`
Histograms and frequency polygons



`geom_jitter()`
Jittered points



`geom_crossbar()` `geom_errorbar()` `geom_linerange()` `geom_pointrange()`
Vertical intervals: lines, crossbars & errorbars



`geom_map()`



Polygons from a reference map

`geom_path()` `geom_line()` `geom_step()`



Connect observations

`geom_point()`



Points

`geom_polygon()`



Polygons

`geom_qq_line()` `stat_qq_line()` `geom_qq()` `stat_qq()`



A quantile-quantile plot

`geom_quantile()` `stat_quantile()`



Quantile regression

`geom_ribbon()` `geom_area()` `stat_align()`



Ribbons and area plots

`geom_rug()`



Rug plots in the margins

`geom_segment()` `geom_curve()`



Line segments and curves

`geom_smooth()` `stat_smooth()`



Smoothed conditional means

`geom_spoke()`



Line segments parameterised by location, direction and distance

`geom_label()` `geom_text()`



Text

`geom_raster()` `geom_rect()` `geom_tile()`



Rectangles

`geom_violin()` `stat_ydensity()`




Violin plot

`coord_sf()` `geom_sf()` `geom_sf_label()` `geom_sf_text()` `stat_sf()`



Visualise sf objects

Next time: Styling



vs

